

Style Transfer

一张是内容图像，另一张是风格图像。

将一个图像中的风格应用在另一图像之上，即风格迁移。

比如说一张相机拍的风光作为内容，梵高的油画作为风格

将风景照片转化成梵高油画风便叫做风格迁移。

Code

```
import torch
import torchvision
from torch import nn
from PIL import Image
from d2l import torch as d2l
import matplotlib.pyplot as plt
from torchvision.models import VGG19_Weights

d2l.set_figsize()
content_img = Image.open('Images/rainier.png').convert('RGB')
style_img = Image.open('Images/autumn.png').convert('RGB')
plt.imshow(content_img)
plt.show()
plt.imshow(style_img)
plt.show()

rgb_mean = torch.tensor([0.485, 0.456, 0.406])
rgb_std = torch.tensor([0.229, 0.224, 0.225])

def preprocess(img, image_shape):
    transforms = torchvision.transforms.Compose([
        torchvision.transforms.Resize(image_shape),
        torchvision.transforms.ToTensor(),
        torchvision.transforms.Normalize(mean=rgb_mean, std=rgb_std)
    ])
    return transforms(img).unsqueeze(0)

# 后处理函数，将tensor转化为图片
def postprocess(img):
    img = img[0].to(rgb_std.device) # img is a tensor
    # permute B C H W -> B H W C
    img = torch.clamp(img.permute(1,2,0) * rgb_std + rgb_mean, 0, 1)
    # B H W C -> B C H W , Tensor -> PIL
    return torchvision.transforms.ToPILImage()(img.permute(2,0,1))

# pretrained=True已经启用，所以使用weights
pretrained_net = torchvision.models.vgg19(weights=VGG19_Weights.IMAGENET1K_V1)

style_layers, content_layers = [0, 5, 10, 19, 28], [25]

net = nn.Sequential(*[pretrained_net.features[i]
                      for i in range(max(content_layers + style_layers) + 1)])

def extract_features(X, content_layers, style_layers):
```

```

contents = []
styles = []
# 对于网络中的每一层
for i in range(len(net)):
    # 在该层上运行输入x以提取特征
    x = net[i](x) # 每一层抽特征
    if i in style_layers: # 如果该层为样式层，则返回样式
        styles.append(x)
    if i in content_layers: # 如果该层为内容层，则返回内容
        contents.append(x)
return contents, styles

def get_contents(image_shape, device):
    content_x = preprocess(content_img, image_shape).to(device)
    # 从内容图像中提取内容特征
    content_Y, _ = extract_features(content_x, content_layers, style_layers)
    return content_x, content_Y

def get_styles(image_shape, device):
    style_x = preprocess(style_img, image_shape).to(device)
    # 从样式图像中提取样式特征
    _, styles_Y = extract_features(style_x, content_layers, style_layers)
    return style_x, styles_Y

def content_loss(Y_hat, Y):
    return torch.square(Y_hat - Y.detach()).mean()

def gram(x):
    # 计算通道数和特征数
    num_channels, n = x.shape[1], x.numel() // x.shape[1]
    # 将输入矩阵reshape为(通道数, 特征数)的格式
    x = x.reshape((num_channels, n))
    return torch.matmul(x, x.T) / (num_channels * n)

def style_loss(Y_hat, gram_Y):
    return torch.square(gram(Y_hat) - gram_Y.detach()).mean()

def tv_loss(Y_hat):
    return 0.5 * (torch.abs(Y_hat[:, :, 1:, :] - Y_hat[:, :, :-1, :]).mean() +
                  torch.abs(Y_hat[:, :, :, 1:] - Y_hat[:, :, :, :-1]).mean())

# 风格转移的损失函数是内容损失、风格损失和总变化损失的加权和
# 定义内容损失、样式损失和总变差损失的权重
content_weight, style_weight, tv_weight = 1, 1e3, 10

# 总损失函数
def compute_loss(X, contents_Y_hat, styles_Y_hat, contents_Y, styles_Y_gram):
    contents_l = [
        content_loss(Y_hat, Y) * content_weight
        for Y_hat, Y in zip(contents_Y_hat, contents_Y) ]
    styles_l = [
        style_loss(Y_hat, Y) * style_weight
        for Y_hat, Y in zip(styles_Y_hat, styles_Y_gram) ]
    tv_l = tv_loss(X) * tv_weight
    l = sum(10 * styles_l + contents_l + [tv_l])
    return contents_l, styles_l, tv_l, l

# 生成合成图像

```

```

class SynthesizedImage(nn.Module):
    def __init__(self, img_shape, **kwargs):
        super(SynthesizedImage, self).__init__(**kwargs)
        self.weight = nn.Parameter(torch.rand(*img_shape))

    def forward(self):
        return self.weight

# 生成图像和优化器
def get_inits(X, device, lr, styles_Y):
    gen_img = SynthesizedImage(X.shape).to(device)
    gen_img.weight.data.copy_(X.data)
    trainer = torch.optim.Adam(gen_img.parameters(), lr=lr)
    styles_Y_gram = [gram(Y) for Y in styles_Y]
    return gen_img(), styles_Y_gram, trainer

def train(X, contents_Y, styles_Y, device, lr, num_epochs, lr_decay_epoch):
    X, styles_Y_gram, trainer = get_inits(X, device, lr, styles_Y)
    scheduler = torch.optim.lr_scheduler.StepLR(trainer, lr_decay_epoch, 0.8)
    animator = d2l.Animator(xlabel='epoch', ylabel='loss', xlim=[10, num_epochs],
                           legend=['content', 'style', 'TV'], ncols=2, figsize=(7, 2.5))
    for epoch in range(num_epochs):
        trainer.zero_grad()
        contents_Y_hat, styles_Y_hat = extract_features(X, content_layers,
        style_layers)
        contents_l, styles_l, tv_l, l = compute_loss(X, contents_Y_hat, styles_Y_hat,
        contents_Y, styles_Y_gram)
        l.backward()
        trainer.step()
        scheduler.step()
        if (epoch+1) % 10 == 0:
            animator.axes[1].imshow(postprocess(X))
            animator.add(epoch + 1,
                        [float(sum(contents_l)),
                        float(sum(styles_l)),
                        float(tv_l)])

    return X

device, image_shape = d2l.try_gpu(), (300, 450)
net = net.to(device)
content_X, contents_Y = get_contents(image_shape, device)
_, styles_Y = get_styles(image_shape, device)
output = train(content_X, contents_Y, styles_Y, device, 0.3, 500, 50)
plt.show(output)

```